

Bee Diversity Analysis using MeanRarity

Analyzing Bee Diversity in a BACI Study Design

This analysis uses the methods described in Roswell et al. (2021) to analyze bee diversity across multiple sites in a Before-After-Control-Impact (BACI) design.

Setup

First, we'll load the necessary packages and read in our data.

```
# Load required packages

library(MeanRarity) # For calculating Hill diversity metrics
library(iNEXT)      # For coverage-based rarefaction
library(dplyr)       # For data manipulation
library(tidyr)       # For reshaping data
library(ggplot2)     # For visualization
library(lubridate)   # For date handling
library(lme4)        # For mixed-effects models
library(lmerTest)    # For p-values in mixed models
library(emmeans)     # For estimated marginal means

# Read in the data
bee_data <- read.csv("../data/processed/bee_02-18_data.csv")
bee_traits <- read.csv("../data/processed/bee_traits.csv", header = TRUE)
site_data <- read.csv("../data/processed/sites.csv", header = TRUE)
dates_to_drop <- read.csv("../data/processed/02_03_18_datestodrop.csv", header = TRUE)
```

Data Cleaning

Clean up the data, standardizing site names and formatting dates.

```
data.frame': 3020 obs. of 11 variables:
 $ barcode      : int  20020087 20020088 20020089 20020090 20020091 20020092 20020094 20020095 ...
 $ genus_name   : chr   "Andrena" "Nomada" "Nomada" "Eucera" ...
 $ species      : chr   "sp. 6" "sp. 3" "sp. 10" "frater" ...
 $ site         : chr   "Goode" "Goode" "Goode" "Goode" ...
 $ start_date   : Date, format: "2002-03-26" "2002-03-26" ...
 $ type         : chr   "Male" "Female" "Female" "Male" ...
 $ det          : chr   "R. W. Brooks" "R. W. Brooks" "R. W. Brooks" "R. W. Brooks" ...
 $ year         : int   2002 2002 2002 2002 2002 2002 2002 2002 2002 2002 ...
 $ number       : int   1 1 1 1 1 1 1 1 1 1 ...
 $ technique    : chr   "netting" "netting" "netting" "netting" ...
 $ combined_name: chr   "Andrena sp. 6" "Nomada sp. 3" "Nomada sp. 10" "Eucera frater" ...
```

Preparing Data for Diversity Analysis

To calculate Hill diversity metrics, we create a matrix of species abundances per site and year.

```
#l warning = FALSE
#l echo = FALSE

# Create a wide-format abundance matrix for diversity analysis
# Each row is a site-year combination, columns are species
abundance_matrix <- bee_data_clean %>%
  group_by(site, year, combined_name) %>%
  # Sum abundance of each species at each site-year
  summarize(abundance = sum(number), .groups = "drop") %>%
  # Convert to wide format: species as columns
  pivot_wider(
    id_cols = c(site, year),
    names_from = combined_name,
    values_from = abundance,
    values_fill = 0 # Fill missing species with 0
  )

# Safely extract site_year_ids with base R
site_year_ids <- data.frame(
  site = abundance_matrix$site,
```

```

    year = abundance_matrix$year
  )

# Safely create abundance_values matrix
# This avoids issues with special characters in species names
abundance_values <- as.matrix(abundance_matrix[, 3:ncol(abundance_matrix)])

# Check dimensions
cat("Number of site-year combinations:", nrow(site_year_ids), "\n")

```

Number of site-year combinations: 24

```

cat("Number of species columns:", ncol(abundance_values), "\n")

```

Number of species columns: 157

```

# Add metadata about sampling effort
sampling_effort <- bee_data_clean %>%
  group_by(site, year) %>%
  summarise(sampling_events = n_distinct(start_date))

```

`summarise()` has grouped output by 'site'. You can override using the
`.groups` argument.

```

# Add treatment information from site_data
site_metadata <- site_year_ids %>%
  left_join(sampling_effort, by = c("site", "year")) %>%
  left_join(site_data %>% dplyr::select(site, treatment), by = "site")

head(site_metadata)

```

	site	year	sampling_events	treatment
1	Bouverie Preserve	2002	8	Impact
2	Bouverie Preserve	2003	7	Impact
3	Bouverie Preserve	2018	8	Impact
4	Goode	2002	11	Control
5	Goode	2003	8	Control
6	Goode	2018	9	Control

Coverage-based Standardization

Calculates coverage. Coverage estimates what proportion of the total community (including undetected species) is represented in our sample. Our results are pretty good.

```
# Calculate coverage for each site-year combination
site_coverage <- data.frame(site_year_ids,
  coverage = apply(abundance_values, 1, function(x) {
    # Count singletons (species with only 1 individual)
    f1 <- sum(x == 1)
    # Get total number of individuals
    n <- sum(x)
    # Calculate coverage using the formula from Chao and Jost (2012)
    if (f1 == 0) {
      return(1) # Perfect coverage if no singletons
    } else {
      return(1 - (f1/n) * ((n-1)/n))
    }
  }))

# Find the minimum coverage across all site-years (to use as standardization point)
min_coverage <- min(site_coverage$coverage)
cat("Minimum coverage across all site-years:", min_coverage, "\n")
```

Minimum coverage across all site-years: 0.7696

```
# Display coverage by site and year
site_coverage %>%
  arrange(coverage) %>%
  knitr::kable(caption = "Sample coverage by site and year")
```

Table 1: Sample coverage by site and year

site	year	coverage
Quintessa	2018	0.7696000
Hudson	2018	0.8310204
Quintessa	2002	0.8630401
Bouverie Preserve	2003	0.8750139
Veterans	2003	0.8889178
Veterans	2018	0.8955464
Saintsbury	2003	0.9043552

site	year	coverage
Stags	2003	0.9109000
Bouverie Preserve	2002	0.9121662
Goode	2018	0.9130481
Saintsbury	2018	0.9222320
Goode	2003	0.9222819
Wappo	2018	0.9240537
Stags	2018	0.9289966
Hudson	2003	0.9337374
Goode	2002	0.9381114
Stags	2002	0.9433975
Quintessa	2003	0.9554158
Bouverie Preserve	2018	0.9555612
Hudson	2002	0.9577470
Veterans	2002	0.9584780
Wappo	2003	0.9735688
Saintsbury	2002	0.9857483
Wappo	2002	0.9897810

Calculating Hill Diversity Metrics

Just some notes on Hill Diversity

Hill diversity unifies different diversity metrics (richness, Shannon, Simpson) into a common framework. The parameter q (or q) determines how much weight is given to rare versus common species.

The MeanRarity package uses parameter q in its code implementation, while their conceptual framework in the paper uses α . The relationship between them is $\alpha = 1 - q$, so:

Species richness: $\alpha = 1$ corresponds to $q = 0$ Hill-Shannon: $\alpha = 0$ corresponds to $q = 1$ Hill-Simpson: $\alpha = -1$ corresponds to $q = 2$

```
# Function to calculate Hill diversity metrics for each site-year
calculate_diversity_metrics <- function() {
  diversity_results <- data.frame()

  for (i in 1:nrow(site_year_ids)) {
    site_id <- site_year_ids$site[i]
    yr <- site_year_ids$year[i]

    # Extract abundance data for this site-year
```

```

site_data_vec <- abundance_values[i, ]
site_data_vec <- site_data_vec[site_data_vec > 0] # Remove zeros

# Calculate raw Hill diversity using the rarity function from MeanRarity
richness_raw <- rarity(site_data_vec, q = 0)
hill_shannon_raw <- rarity(site_data_vec, q = 1)
hill_simpson_raw <- rarity(site_data_vec, q = 2)

# Use iNEXT to estimate diversity at standardized coverage
inext_result <- estimateD(site_data_vec, datatype = "abundance",
                          base = "coverage", level = min_coverage)

# The issue is in extracting values from inext_result
# Print the structure to understand how to access the values
# print(str(inext_result))

# Get the values in the correct way
# iNEXT results are sorted by order (0, 1, 2)
richness_std <- inext_result$qD[1] # First row is q=0 (richness)
hill_shannon_std <- inext_result$qD[2] # Second row is q=1 (Shannon)
hill_simpson_std <- inext_result$qD[3] # Third row is q=2 (Simpson)

# Combine with existing metadata
result_row <- data.frame(
  site = site_id,
  year = yr,
  raw_richness = richness_raw,
  raw_hill_shannon = hill_shannon_raw,
  raw_hill_simpson = hill_simpson_raw,
  std_richness = richness_std,
  std_hill_shannon = hill_shannon_std,
  std_hill_simpson = hill_simpson_std,
  coverage = site_coverage$coverage[i]
)

# Append to the result data frame
diversity_results <- rbind(diversity_results, result_row)
}

# Add metadata about treatment and sampling effort
diversity_results <- diversity_results %>%
  left_join(site_metadata, by = c("site", "year"))

```

```

    return(diversity_results)
}

# Calculate diversity metrics
diversity_results <- calculate_diversity_metrics()

# Create factors for time period and ensure proper ordering
diversity_results$period <- ifelse(diversity_results$year == "2018", "After", "Before")
diversity_results$period <- factor(diversity_results$period, levels = c("Before", "After"))
diversity_results$treatment <- factor(diversity_results$treatment, levels = c("Control", "Imp"))
diversity_results$year <- factor(diversity_results$year, levels = c("2002", "2003", "2018"))

```


Visualizing Diversity Metrics

Visualization of the diversity results to understand the patterns across sites, years, and treatments.

```


# Visualize Hill-Shannon diversity (standardized by coverage)
ggplot(diversity_results, aes(x = year, y = std_hill_shannon, color = treatment, group = site)) +
  geom_point(size = 3) +
  geom_line(linewidth = 1) +
  labs(title = "Hill-Shannon Diversity (Coverage-Standardized)",
       x = "Year",
       y = "Hill-Shannon Diversity (q=1)") +
  theme_minimal()

```



```
# Compare all three diversity metrics
diversity_long <- diversity_results %>%
  pivot_longer(cols = c(std_richness, std_hill_shannon, std_hill_simpson),
    names_to = "metric",
    values_to = "value") %>%
  mutate(metric = factor(metric,
    levels = c("std_richness", "std_hill_shannon", "std_hill_simpson"),
    labels = c("Species Richness (q=0)",
      "Hill-Shannon (q=1)",
      "Hill-Simpson (q=2)")))

ggplot(diversity_long, aes(x = year, y = value, color = treatment, group = site)) +
  facet_wrap(~ metric, scales = "free_y") +
  geom_point(size = 2) +
  geom_line(linewidth = 1) +
  labs(title = "Coverage-Standardized Hill Diversity Metrics",
    x = "Year",
    y = "Diversity Value") +
  theme_minimal()
```

Diversity Profiles

Diversity profiles show how diversity changes with the Hill exponent , giving a more complete picture of community structure. This is just done for 2018. Could do for all years.

```
# Create diversity profiles for each site in the final year (2018)
# Filter to just 2018 data
data_2018 <- abundance_matrix %>%
  filter(year == "2018")
site_ids_2018 <- data_2018$site

# Safely extract abundance data for 2018 sites
# Using the same approach we used above
abundance_2018 <- as.matrix(data_2018[, 3:ncol(data_2018)])

# Get treatment information
treatments_2018 <- site_metadata %>%
  filter(year == "2018") %>%
  dplyr::select(site, treatment)

# Set up the plot
```

```

par(mfrow = c(1, 1))
plot(NULL, xlim = c(-1, 1), ylim = c(0, max(diversity_results$std_richness) * 1.1),
     xlab = "Hill exponent ( )", ylab = "Hill Diversity",
     main = "Diversity Profiles (2018)")

# Define colors for treatments
colors <- c("Control" = "blue", "Impact" = "red")

# Calculate and plot diversity profiles
for (i in 1:nrow(abundance_2018)) {
  site_id <- site_ids_2018[i]
  site_abund <- abundance_2018[i, ]
  site_abund <- site_abund[site_abund > 0] # Remove zeros

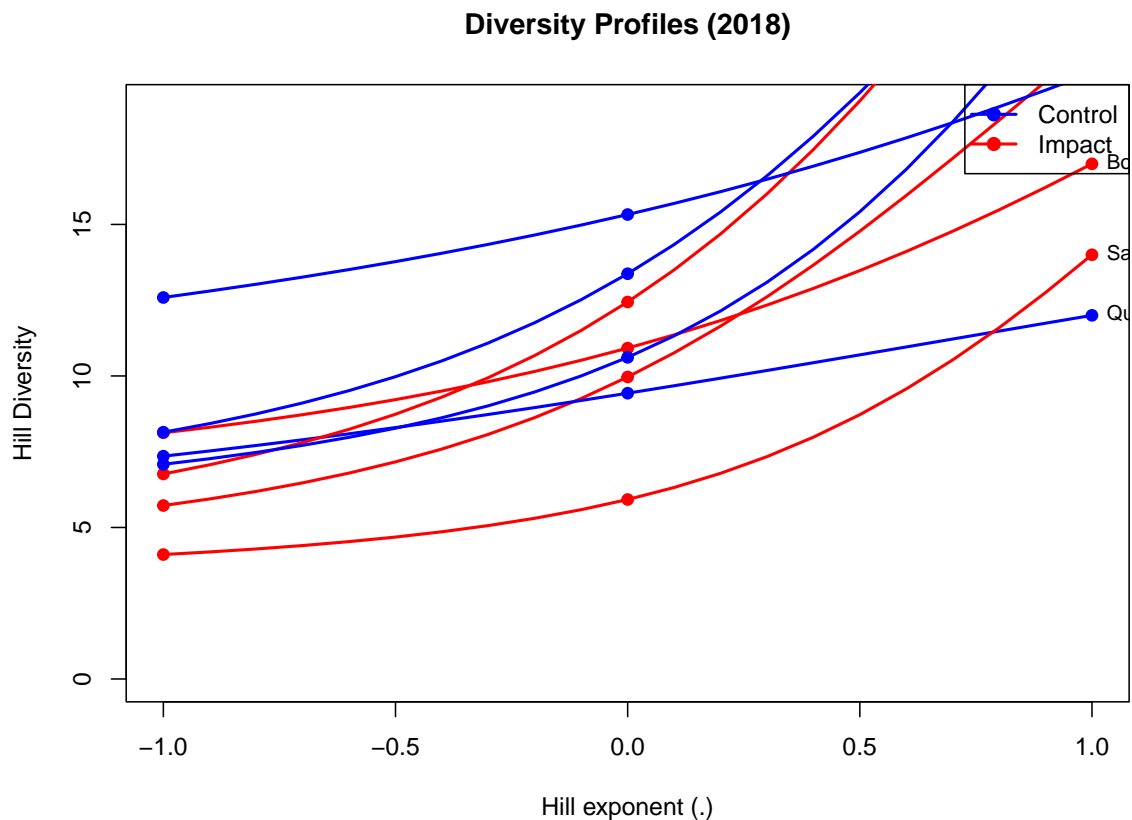
  # Get treatment for this site
  site_treatment <- treatments_2018$treatment[treatments_2018$site == site_id]

  # Calculate diversity for a range of exponents
  exponents <- seq(-1, 1, by = 0.1)
  profile_values <- sapply(exponents, function(ell) rarity(site_abund, q = 1 - ell))

  # Plot the profile
  lines(exponents, profile_values, col = colors[site_treatment], lwd = 2)
  points(c(-1, 0, 1), profile_values[c(1, 11, 21)],
        col = colors[site_treatment], pch = 19)
  text(1, profile_values[21], site_id, pos = 4, cex = 0.8)
}

# Add a legend
legend("topright", legend = names(colors), col = colors, lwd = 2, pch = 19)

```



Statistical Analysis

Linear mixed-effects models with a BACI design. Dependent variable is Hill-Shannon diversity but also analyze richness and Simpson diversity. Repeated the model with and without taking into account sampling effort. AICs for models with and without taking sampling effort into account are virtually identical

```
# Model for Hill-Shannon diversity
model_hill_shannon <- lmer(std_hill_shannon ~ year * treatment + (1 | site),
                           data = diversity_results)
summary(model_hill_shannon)
```

```
Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]
Formula: std_hill_shannon ~ year * treatment + (1 | site)
Data: diversity_results
```

REML criterion at convergence: 90.2

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.7334	-0.6492	0.1059	0.7308	1.3771

Random effects:

Groups	Name	Variance	Std.Dev.
site	(Intercept)	0.8421	0.9176
Residual		4.8132	2.1939

Number of obs: 24, groups: site, 8

Fixed effects:

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	6.818	1.189	17.236	5.734	2.31e-05 ***
year2003	1.121	1.551	12.000	0.723	0.484
year2018	2.710	1.551	12.000	1.747	0.106
treatmentImpact	-1.047	1.682	17.236	-0.623	0.542
year2003:treatmentImpact	1.511	2.194	12.000	0.689	0.504
year2018:treatmentImpact	-1.247	2.194	12.000	-0.568	0.580

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	yr2003	yr2018	trtmnI	y2003:
year2003		-0.652			
year2018		-0.652	0.500		
trtmntImpct		-0.707	0.461	0.461	
yr2003:trtI		0.461	-0.707	-0.354	-0.652
yr2018:trtI		0.461	-0.354	-0.707	-0.652
					0.500

```
# Model for Hill-Shannon diversity including sampling events
model_hill_shannon_with_sampling <- lmer(std_hill_shannon ~ year * treatment + sampling_events,
                                          data = diversity_results)
summary(model_hill_shannon_with_sampling)
```

Linear mixed model fit by REML. t-tests use Satterthwaite's method [lmerModLmerTest]

Formula: std_hill_shannon ~ year * treatment + sampling_events + (1 | site)

Data: diversity_results

REML criterion at convergence: 90.1

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.5857	-0.5336	0.1877	0.5641	1.3678

Random effects:

Groups	Name	Variance	Std.Dev.
site	(Intercept)	1.435	1.198
Residual		4.666	2.160

Number of obs: 24, groups: site, 8

Fixed effects:

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	4.8545	3.8724	14.6372	1.254	0.2296
year2003	1.4312	1.6337	12.2661	0.876	0.3978
year2018	3.1237	1.7118	13.2722	1.825	0.0906
treatmentImpact	-0.9957	1.7493	14.7115	-0.569	0.5778
sampling_events	0.2067	0.3863	14.6287	0.535	0.6007
year2003:treatmentImpact	1.6144	2.1688	10.7831	0.744	0.4726
year2018:treatmentImpact	-1.8150	2.4073	13.1567	-0.754	0.4641

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	yr2003	yr2018	trtmnI	smpln_	y2003:
year2003		-0.521				
year2018		-0.604	0.577			
trtmntImpct		-0.278	0.428	0.415		
smplng_vnts		-0.948	0.355	0.451	0.055	
yr2003:trtI		0.054	-0.627	-0.274	-0.610	0.089
yr2018:trtI		0.543	-0.453	-0.765	-0.578	-0.441

```
# Compare models
anova(model_hill_shannon, model_hill_shannon_with_sampling)
```

refitting model(s) with ML (instead of REML)

Data: diversity_results

Models:

model_hill_shannon: std_hill_shannon ~ year * treatment + (1 | site)

```

model_hill_shannon_with_sampling: std_hill_shannon ~ year * treatment + sampling_events + (1
                                npar      AIC      BIC logLik deviance Chisq Df
model_hill_shannon              8 118.29 127.72 -51.146   102.29
model_hill_shannon_with_sampling 9 120.02 130.63 -51.012   102.02 0.268  1
                                Pr(>Chisq)

model_hill_shannon
model_hill_shannon_with_sampling      0.6047

```

```

# Extract p-values
p_values <- summary(model_hill_shannon)$coefficients[, "Pr(>|t|)"]
print(p_values)

```

(Intercept)	year2003	year2018
2.307004e-05	4.837059e-01	1.061402e-01
treatmentImpact	year2003:treatmentImpact	year2018:treatmentImpact
5.415425e-01	5.040872e-01	5.803520e-01

```

# Calculate estimated marginal means
emm_hill_shannon <- emmeans(model_hill_shannon, ~ year * treatment)
pairs(emm_hill_shannon)

```

contrast	estimate	SE	df	t.ratio	p.value
year2002 Control - year2003 Control	-1.121	1.55	12.0	-0.723	0.9752
year2002 Control - year2018 Control	-2.710	1.55	12.0	-1.747	0.5295
year2002 Control - year2002 Impact	1.047	1.68	17.2	0.623	0.9877
year2002 Control - year2003 Impact	-1.585	1.68	17.2	-0.942	0.9296
year2002 Control - year2018 Impact	-0.416	1.68	17.2	-0.248	0.9998
year2003 Control - year2018 Control	-1.589	1.55	12.0	-1.024	0.9009
year2003 Control - year2002 Impact	2.168	1.68	17.2	1.290	0.7866
year2003 Control - year2003 Impact	-0.464	1.68	17.2	-0.276	0.9997
year2003 Control - year2018 Impact	0.705	1.68	17.2	0.419	0.9980
year2018 Control - year2002 Impact	3.758	1.68	17.2	2.235	0.2722
year2018 Control - year2003 Impact	1.125	1.68	17.2	0.669	0.9831
year2018 Control - year2018 Impact	2.294	1.68	17.2	1.364	0.7466
year2002 Impact - year2003 Impact	-2.632	1.55	12.0	-1.697	0.5580
year2002 Impact - year2018 Impact	-1.464	1.55	12.0	-0.943	0.9271
year2003 Impact - year2018 Impact	1.169	1.55	12.0	0.753	0.9705

Degrees-of-freedom method: kenward-roger

P value adjustment: tukey method for comparing a family of 6 estimates

```
# Calculate estimated marginal means accounting for sampling events
# Using emmeans for a specific level of sampling_events (mean value)
mean_sampling <- mean(diversity_results$sampling_events)
emm_hill_shannon_sampling <- emmeans(model_hill_shannon_with_sampling,
                                     ~ year * treatment,
                                     at = list(sampling_events = mean_sampling))
pairs(emm_hill_shannon_sampling)
```

contrast	estimate	SE	df	t.ratio	p.value
year2002 Control - year2003 Control	-1.431	1.66	12.7	-0.861	0.9494
year2002 Control - year2018 Control	-3.124	1.76	13.7	-1.774	0.5114
year2002 Control - year2002 Impact	0.996	1.75	15.0	0.569	0.9917
year2002 Control - year2003 Impact	-2.050	2.01	15.0	-1.022	0.9034
year2002 Control - year2018 Impact	-0.313	1.76	15.0	-0.178	1.0000
year2003 Control - year2018 Control	-1.692	1.54	11.4	-1.097	0.8732
year2003 Control - year2002 Impact	2.427	1.83	15.0	1.326	0.7672
year2003 Control - year2003 Impact	-0.619	1.78	15.0	-0.348	0.9992
year2003 Control - year2018 Impact	1.118	1.95	15.0	0.572	0.9915
year2018 Control - year2002 Impact	4.119	1.91	15.0	2.159	0.3109
year2018 Control - year2003 Impact	1.074	1.75	15.0	0.614	0.9883
year2018 Control - year2018 Impact	2.811	2.06	15.0	1.363	0.7470
year2002 Impact - year2003 Impact	-3.046	1.76	13.7	-1.729	0.5367
year2002 Impact - year2018 Impact	-1.309	1.56	11.6	-0.838	0.9541
year2003 Impact - year2018 Impact	1.737	1.95	15.0	0.893	0.9425

Degrees-of-freedom method: kenward-roger

P value adjustment: tukey method for comparing a family of 6 estimates

```
# Also analyze richness and Simpson diversity for comparison
model_richness <- lmer(std_richness ~ year * treatment + (1 | site),
                      data = diversity_results)
summary(model_richness)
```

Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]

Formula: std_richness ~ year * treatment + (1 | site)

Data: diversity_results

REML criterion at convergence: 105.5

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.70784	-0.50146	-0.02167	0.58058	1.36244

Random effects:

Groups	Name	Variance	Std.Dev.
site	(Intercept)	0.2659	0.5156
Residual		12.6736	3.5600

Number of obs: 24, groups: site, 8

Fixed effects:

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	9.809	1.799	17.985	5.454	3.53e-05 ***
year2003	1.156	2.517	12.000	0.459	0.654
year2018	2.475	2.517	12.000	0.983	0.345
treatmentImpact	-2.727	2.544	17.985	-1.072	0.298
year2003:treatmentImpact	4.236	3.560	12.000	1.190	0.257
year2018:treatmentImpact	1.088	3.560	12.000	0.306	0.765

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	yr2003	yr2018	trtmnI	y2003:
year2003	-0.700				
year2018	-0.700	0.500			
trtmntImpct	-0.707	0.495	0.495		
yr2003:trtI	0.495	-0.707	-0.354	-0.700	
yr2018:trtI	0.495	-0.354	-0.707	-0.700	0.500

```
# Update richness model to include sampling events
model_richness_with_sampling <- lmer(std_richness ~ year * treatment + sampling_events + (1
                                   data = diversity_results)
summary(model_richness_with_sampling)
```

Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]

Formula: std_richness ~ year * treatment + sampling_events + (1 | site)

Data: diversity_results

REML criterion at convergence: 104.8

Scaled residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

-1.69154 -0.37122 0.01469 0.51107 1.40249

Random effects:

Groups	Name	Variance	Std.Dev.
site	(Intercept)	1.374	1.172
Residual		12.426	3.525

Number of obs: 24, groups: site, 8

Fixed effects:

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	7.9015	5.7339	11.9698	1.378	0.193
year2003	1.4571	2.6356	11.6763	0.553	0.591
year2018	2.8767	2.7417	12.8212	1.049	0.313
treatmentImpact	-2.6766	2.6307	16.2499	-1.017	0.324
sampling_events	0.2008	0.5710	11.4175	0.352	0.732
year2003:treatmentImpact	4.3367	3.5365	10.0903	1.226	0.248
year2018:treatmentImpact	0.5362	3.8590	12.6869	0.139	0.892

Correlation of Fixed Effects:

	(Intr)	yr2003	yr2018	trtmnI	smpln_	y2003:
year2003	-0.513					
year2018	-0.592	0.565				
trtmntImpct	-0.280	0.466	0.453			
smplng_vnts	-0.946	0.325	0.417	0.054		
yr2003:trtI	0.077	-0.640	-0.287	-0.663	0.081	
yr2018:trtI	0.525	-0.438	-0.757	-0.634	-0.407	0.422

```
# Compare richness models
anova(model_richness, model_richness_with_sampling)
```

refitting model(s) with ML (instead of REML)

Data: diversity_results

Models:

model_richness: std_richness ~ year * treatment + (1 | site)

model_richness_with_sampling: std_richness ~ year * treatment + sampling_events + (1 | site)

	npar	AIC	BIC	logLik	deviance	Chisq	Df
model_richness	8	138.64	148.07	-61.321	122.64		
model_richness_with_sampling	9	140.58	151.18	-61.290	122.58	0.0616	1

Pr(>Chisq)

model_richness

model_richness_with_sampling 0.804

```
model_simpson <- lmer(std_hill_simpson ~ year * treatment + (1 | site),
                      data = diversity_results)
summary(model_simpson)
```

```
Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]
Formula: std_hill_simpson ~ year * treatment + (1 | site)
Data: diversity_results
```

REML criterion at convergence: 79.9

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.73870	-0.59927	0.06513	0.63334	1.50122

Random effects:

Groups	Name	Variance	Std.Dev.
site	(Intercept)	0.619	0.7867
Residual		2.612	1.6162

Number of obs: 24, groups: site, 8

Fixed effects:

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	5.0304	0.8987	16.7691	5.597	3.37e-05 ***
year2003	0.7827	1.1428	12.0000	0.685	0.5064
year2018	2.5056	1.1428	12.0000	2.193	0.0488 *
treatmentImpact	-0.2248	1.2710	16.7691	-0.177	0.8618
year2003:treatmentImpact	0.4380	1.6162	12.0000	0.271	0.7910
year2018:treatmentImpact	-2.0831	1.6162	12.0000	-1.289	0.2217

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	yr2003	yr2018	trtmntI	y2003:
year2003	-0.636				
year2018	-0.636	0.500			
trtmntImpct	-0.707	0.450	0.450		
yr2003:trtI	0.450	-0.707	-0.354	-0.636	
yr2018:trtI	0.450	-0.354	-0.707	-0.636	0.500

```
# Update Simpson model to include sampling events
model_simpson_with_sampling <- lmer(std_hill_simpson ~ year * treatment + sampling_events +
                                     data = diversity_results)
summary(model_simpson_with_sampling)
```

Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]

Formula: std_hill_simpson ~ year * treatment + sampling_events + (1 |
site)

Data: diversity_results

REML criterion at convergence: 80.4

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.5794	-0.5174	0.0958	0.6229	1.4835

Random effects:

Groups	Name	Variance	Std.Dev.
site	(Intercept)	0.8641	0.9296
	Residual	2.6060	1.6143

Number of obs: 24, groups: site, 8

Fixed effects:

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	3.8388	2.9201	15.0585	1.315	0.208
year2003	0.9708	1.2223	12.6771	0.794	0.442
year2018	2.7565	1.2816	13.6110	2.151	0.050 *
treatmentImpact	-0.1934	1.3192	14.7667	-0.147	0.885
sampling_events	0.1254	0.2913	15.0909	0.431	0.673
year2003:treatmentImpact	0.5007	1.6209	11.2709	0.309	0.763
year2018:treatmentImpact	-2.4281	1.8022	13.5045	-1.347	0.200

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	yr2003	yr2018	trtmnI	smpln_	y2003:
year2003		-0.521				
year2018		-0.605	0.578			
trtmntImpct		-0.278	0.424	0.410		
smplng_vnts		-0.948	0.358	0.455	0.055	
yr2003:trtI		0.052	-0.626	-0.273	-0.604	0.090

```
yr2018:trtI  0.545 -0.455 -0.766 -0.573 -0.445  0.406
```

```
# Compare Simpson models
anova(model_simpson, model_simpson_with_sampling)
```

```
refitting model(s) with ML (instead of REML)
```

```
Data: diversity_results
```

```
Models:
```

```
model_simpson: std_hill_simpson ~ year * treatment + (1 | site)
```

```
model_simpson_with_sampling: std_hill_simpson ~ year * treatment + sampling_events + (1 | si
```

	npar	AIC	BIC	logLik	deviance	Chisq	Df
model_simpson	8	104.54	113.97	-44.272	88.543		
model_simpson_with_sampling	9	106.34	116.94	-44.170	88.340	0.2036	1

Pr(>Chisq)

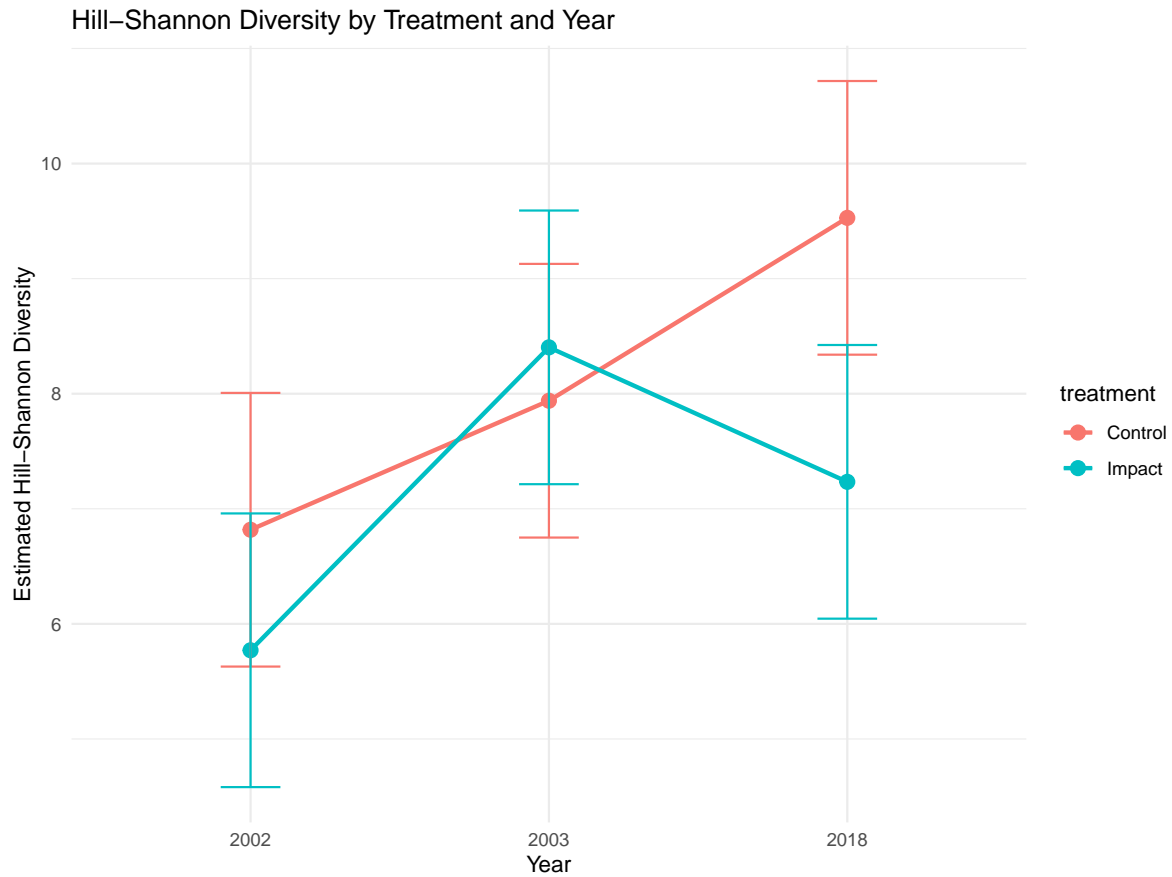
```
model_simpson
```

```
model_simpson_with_sampling 0.6518
```

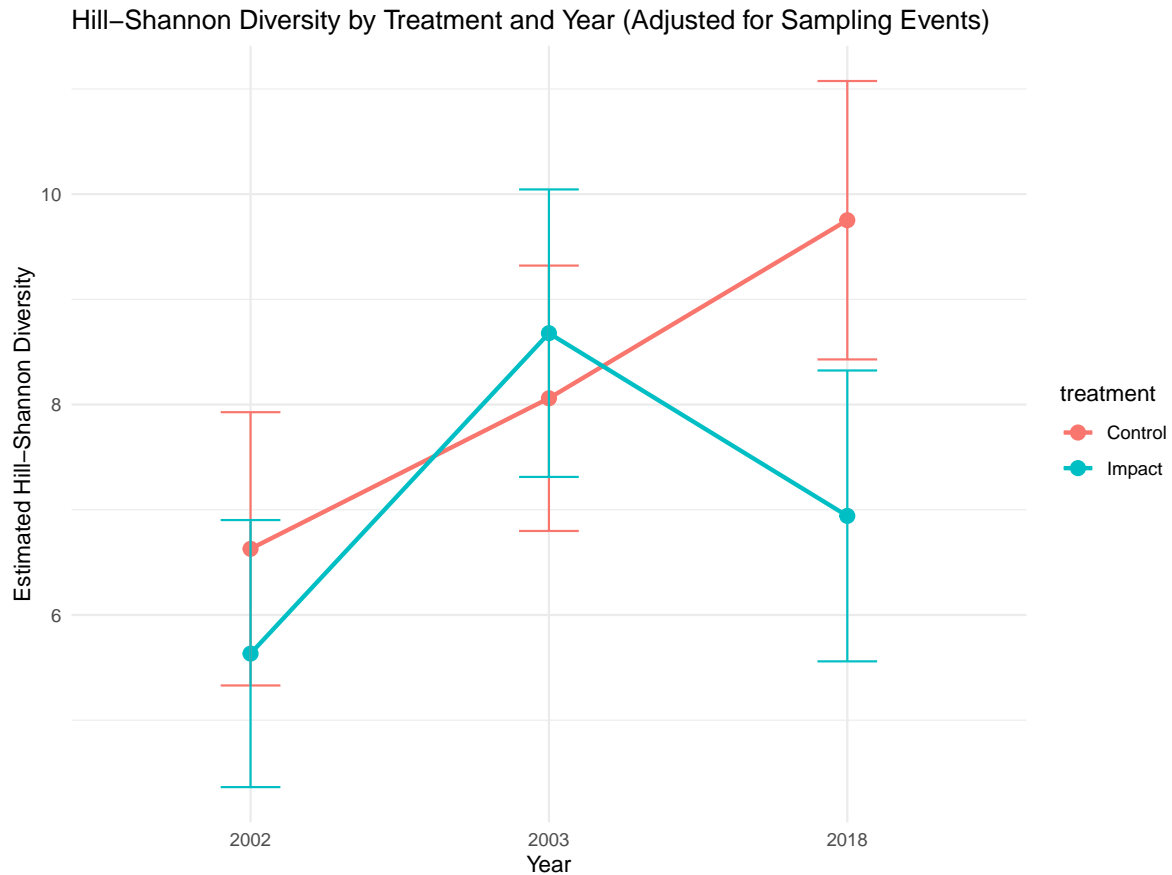
```
# Plot the estimated means with error bars
```

```
emm_data <- as.data.frame(emm_hill_shannon)
```

```
ggplot(emm_data, aes(x = year, y = emmean, color = treatment, group = treatment)) +
  geom_point(size = 3) +
  geom_line(linewidth = 1) +
  geom_errorbar(aes(ymin = emmean - SE, ymax = emmean + SE), width = 0.2) +
  labs(title = "Hill-Shannon Diversity by Treatment and Year",
       x = "Year",
       y = "Estimated Hill-Shannon Diversity") +
  theme_minimal()
```



```
# Plot the estimated means with error bars, adjusted for sampling events
emm_data_sampling <- as.data.frame(emm_hill_shannon_sampling)
ggplot(emm_data_sampling, aes(x = year, y = emmean, color = treatment, group = treatment)) +
  geom_point(size = 3) +
  geom_line(linewidth = 1) +
  geom_errorbar(aes(ymin = emmean - SE, ymax = emmean + SE), width = 0.2) +
  labs(title = "Hill-Shannon Diversity by Treatment and Year (Adjusted for Sampling Events)"
        x = "Year",
        y = "Estimated Hill-Shannon Diversity") +
  theme_minimal()
```



Analysis of Raw vs. Coverage-standardized Results

Compares the raw diversity metrics with the coverage-standardized metrics to see how standardization affects our conclusions.

Also does some plotting of effect of `sampling_events`.

The shape of the curves are similar although estimated species of the Hill Shannon estimate is lower.

```
# Combine raw and standardized metrics for comparison
diversity_compare <- diversity_results %>%
  dplyr::select(site, year, treatment, period,
    raw_hill_shannon, std_hill_shannon,
    sampling_events, coverage) %>%
  pivot_longer(
    cols = c(raw_hill_shannon, std_hill_shannon),
```


```

names_to = "metric_type",
values_to = "hill_shannon"
) %>%
mutate(metric_type = factor(metric_type,
                             levels = c("raw_hill_shannon", "std_hill_shannon"),
                             labels = c("Raw Hill-Shannon", "Coverage-standardized Hill-Shannon"))

# Plot the comparison
ggplot(diversity_compare, aes(x = year, y = hill_shannon, linetype = treatment, color = site)) +
  facet_wrap(~ metric_type) +
  geom_point(size = 3) +
  geom_line(size = 1) +
  labs(title = "Raw vs. Coverage-standardized Hill-Shannon Diversity",
       x = "Year",
       y = "Hill-Shannon Diversity") +
  theme_minimal()

```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
 i Please use `linewidth` instead.



```
# Check correlation between sampling events and raw diversity
cor_sampling_raw <- cor(diversity_results$sampling_events, diversity_results$raw_hill_shannon)
cat("Correlation between sampling events and raw Hill-Shannon diversity:", cor_sampling_raw,
```

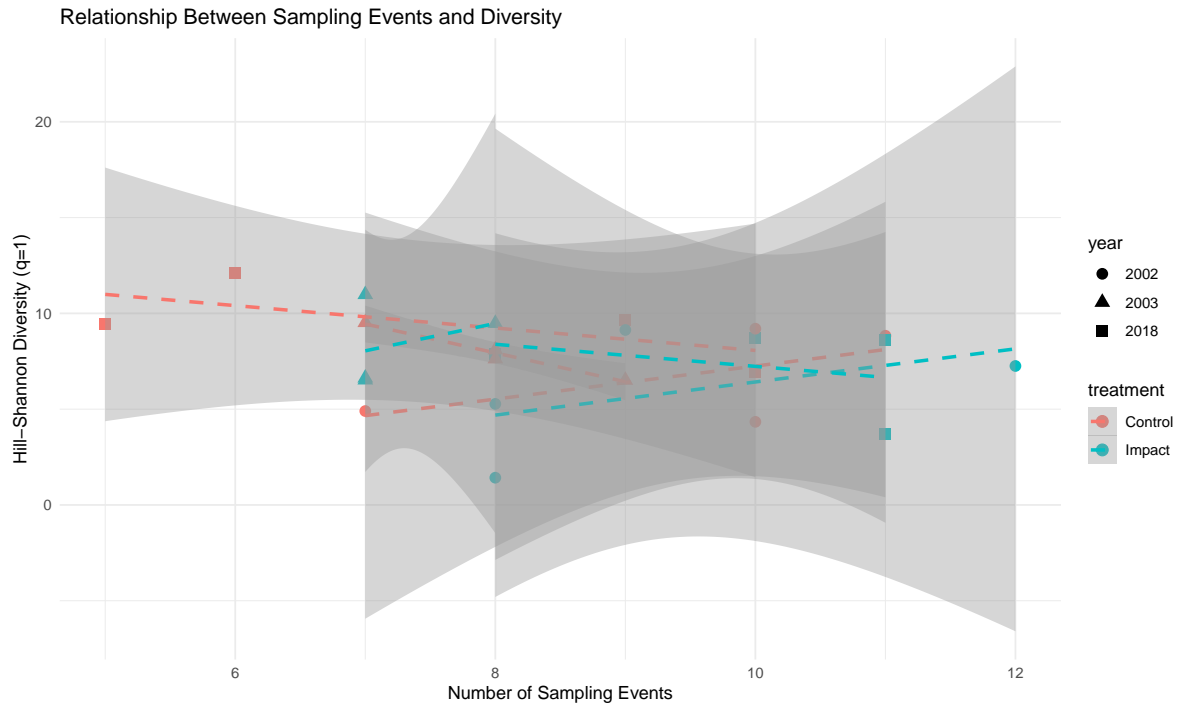
Correlation between sampling events and raw Hill-Shannon diversity: 0.05416637

```
# Check correlation between sampling events and standardized diversity
cor_sampling_std <- cor(diversity_results$sampling_events, diversity_results$std_hill_shannon)
cat("Correlation between sampling events and standardized Hill-Shannon diversity:", cor_samp
```

Correlation between sampling events and standardized Hill-Shannon diversity: -0.2178351


```
# Visualize diversity by sampling events
ggplot(diversity_results, aes(x = sampling_events, y = std_hill_shannon, color = treatment, shape = treatment)) +
  geom_point(size = 3) +
  geom_smooth(method = "lm", se = TRUE, linetype = "dashed") +
  labs(title = "Relationship Between Sampling Events and Diversity",
       x = "Number of Sampling Events",
       y = "Hill-Shannon Diversity (q=1)") +
  theme_minimal()
```

`geom_smooth()` using formula = 'y ~ x'



```
# Add facets to see the relationship by year
ggplot(diversity_results, aes(x = sampling_events, y = std_hill_shannon, color = treatment)) +
  geom_point(size = 3) +
  geom_smooth(method = "lm", se = TRUE, linetype = "dashed") +
  facet_wrap(~ year) +
  labs(title = "Relationship Between Sampling Events and Diversity by Year",
       x = "Number of Sampling Events",
       y = "Hill-Shannon Diversity (q=1)") +
  theme_minimal()
```

`geom_smooth()` using formula = 'y ~ x'



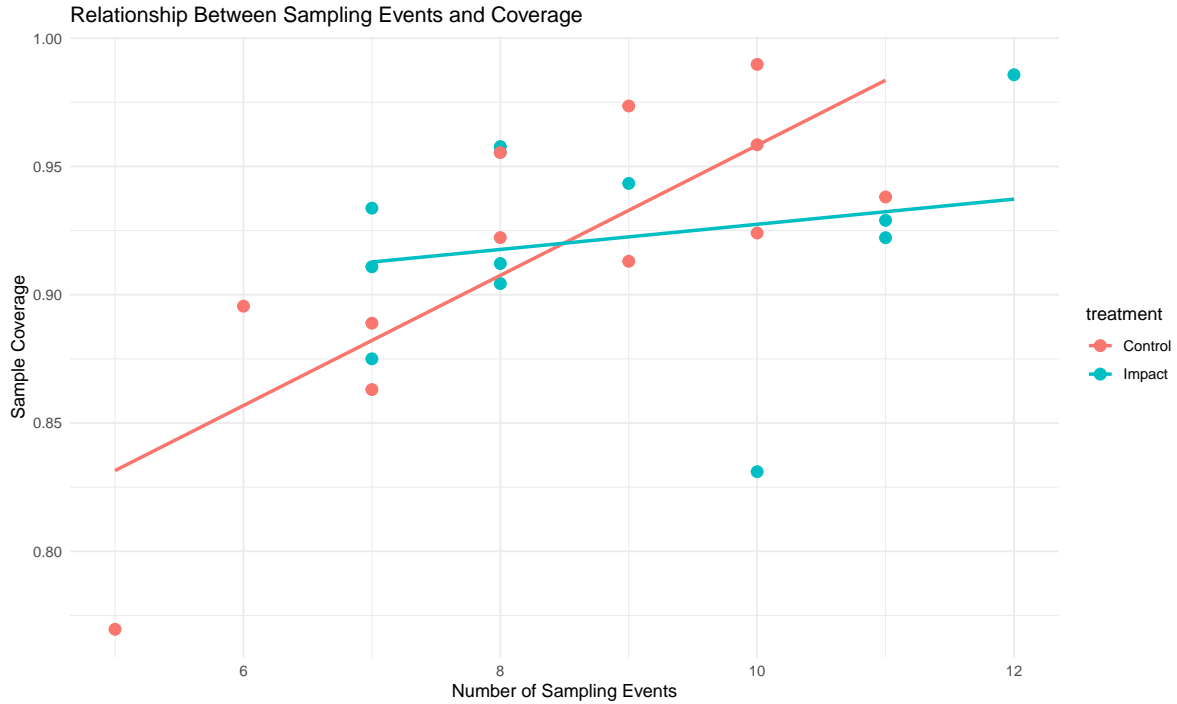
```
# Compare correlations across all diversity metrics
cor_data <- data.frame(
  metric = c("Raw Richness", "Raw Hill-Shannon", "Raw Hill-Simpson",
            "Std Richness", "Std Hill-Shannon", "Std Hill-Simpson"),
  correlation = c(
    cor(diversity_results$sampling_events, diversity_results$raw_richness),
    cor(diversity_results$sampling_events, diversity_results$raw_hill_shannon),
    cor(diversity_results$sampling_events, diversity_results$raw_hill_simpson),
    cor(diversity_results$sampling_events, diversity_results$std_richness),
    cor(diversity_results$sampling_events, diversity_results$std_hill_shannon),
    cor(diversity_results$sampling_events, diversity_results$std_hill_simpson)
  ),
  metric_type = rep(c("Raw", "Standardized"), each = 3)
)

# Plot correlations
ggplot(cor_data, aes(x = metric, y = correlation, fill = metric_type)) +
  geom_col() +
  coord_flip() +
  labs(title = "Correlation Between Sampling Events and Diversity Metrics",
       y = "Pearson Correlation Coefficient") +
  theme_minimal()
```



```
# Visualize relationship between sampling events and coverage
ggplot(diversity_results, aes(x = sampling_events, y = coverage, color = treatment)) +
  geom_point(size = 3) +
  geom_smooth(method = "lm", se = FALSE) +
  labs(title = "Relationship Between Sampling Events and Coverage",
       x = "Number of Sampling Events",
       y = "Sample Coverage") +
  theme_minimal()
```

`geom_smooth()` using formula = 'y ~ x'



Notes and resultss.

Hill Diversity Metrics: The three Hill diversity metrics (richness, Shannon, and Simpson) provide different perspectives on community diversity:

- Species richness ($=1, q=0$) gives equal weight to all species and is most sensitive to rare species
- Hill-Shannon diversity ($=0, q=1$) provides a balanced view of diversity
- Hill-Simpson diversity ($=-1, q=2$) emphasizes common species

All Hill diversity metrics are expressed in “effective number of species,” making them more intuitive to interpret than traditional indices.

1. **Coverage Standardization:** The minimum coverage across our samples was 0.769, which means we’re comparing equally complete samples from each community.

2. **BACI Analysis:**

- None of the models are significant
- We have no interaction terms that are significant
- Using the Simpson diversity which emphasizes common diversity, there is a difference between 2002 and 2018

3. **Sampling Effort Analysis:** Our analysis of sampling effort shows:

- The relationship between the number of sampling events and both raw and coverage-standardized diversity metrics
- How well coverage standardization accounts for differences in sampling effort
- Whether accounting for sampling events in our statistical models changes our conclusions about treatment effects
-

References

Roswell, M., Dushoff, J., & Winfree, R. (2021). A conceptual guide to measuring species diversity. *Oikos*, 130(3), 321-338.